

Joins in SQL Server



Klaus Aschenbrenner


Microsoft Certified Master SQL Server 2008

www.SQLpassion.at

Twitter: @Aschenbrenner

About me

- CEO & Founder SQLpassion
- International Speaker, Blogger, Author
- SQL Server 2008 MCM
- "Pro SQL Server 2008 Service Broker"
- Twitter: @Aschenbrenner
- SQLpassion Academy
 - <http://www.SQLpassion.at/academy>
 - Free Newsletter, Training Videos



Agenda

- Physical Joins
- Logical Joins

Agenda


- Physical Joins
- Logical Joins

Query Trees

```
SELECT
    A.ID,
    B.ID
FROM A
JOIN B ON A.ID = B.ID
WHERE A.X = <Value>
```


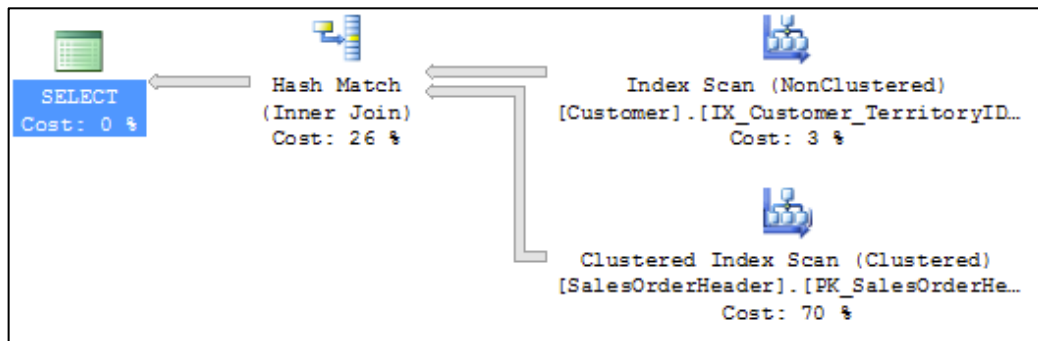
Query Trees

```
SELECT  
    A.ID,  
    B.ID  
FROM A  
JOIN B ON A.ID = B.ID  
WHERE A.X = <Value>
```



Query Trees

```
SELECT
    A.ID,
    B.ID
FROM A
JOIN B ON A.ID = B.ID
WHERE A.X = <Value>
```



Nested Loop

- "For each Row" operator
- Takes output from one step and executes another operation "for each" output row
- Outer Loop, Inner Loop
- Only join type that supports inequality Predicates

Merge Join

- Needs at least one equijoin predicate
- Used when joined columns are indexed (sorted)
- Otherwise (expensive) sorting is needed
 - Plan may include explicit sort

Hash Join

- Needs at least one equijoin predicate
- Hashes values of join columns from one side (smaller table)
 - Based on Statistics
- Probes them with the join columns of the other side (larger table)
- Stop and Go for the Probe Phase
- Needs memory grants to build the hash table
 - If the memory grants are exceeded, the Hash Join is spilled to TempDb
 - Performance decreases!

Demo

Physical Joins

Agenda

- Physical Joins
- Logical Joins

Joins

- Operates on 2 input tables
- 3 fundamental types
 - CROSS JOIN
 - INNER JOIN
 - OUTER JOIN
- Differ in how they apply their logical query processing phases

Logical Phases

- Each Join type applies a different set of phases
 - Cartesian Product
 - Filter
 - Add Outer Rows

Cross Joins

- Simplest type of Join
- Implement only one logical query processing phase
 - Cartesian Product
 - Each row from each input is matched with the rows from each other input
 - Result: $m * n$ rows

Demo

Cross Joins

Inner Joins

- Applies 2 logical query processing phases
 - Cartesian Product
 - Filters rows based on a predicate
- Only returns rows where the predicate returns TRUE
 - FALSE, UNKNOWN are not returned
- INNER keyword is optional
- Use the new ANSI SQL-92 syntax to avoid Cross Joins
 - ON clause is needed

Non-Equi Joins

- Join condition involves an non-equality operator
 - <
 - >
 - <=
 - >=
 - <>
- Implemented as a Nested Loop operator

Demo

Inner Joins

Outer Joins

- Apply 3 logical query processing phases
 - Cartesian Product
 - Filter rows based on a predicate
 - Adding Outer Rows
- One table is marked as „preserved“
 - LEFT OUTER JOIN
 - RIGHT OUTER JOIN
 - FULL OUTER JOIN
- OUTER keyword is optional
- Rows from the preserved table are added to the other table
 - With NULL values as placeholders

Demo

Outer Joins

Summary

- Physical Joins
- Logical Joins

SQL Server Performance Tuning Workshop

- June 1 – 5 in London
- Agenda
 - Database Internals
 - Execution Plans
 - Indexing
 - Statistics
 - Locking, Blocking, Deadlocking
 - Performance Troubleshooting
- More Information
 - <http://www.SQLpassion.at/academy/perftuning>
 - 10% Discount!

