



Strategic Sponsor



Gold Sponsors



Silver Sponsors



Technical Partners



Academic Partners



Wyższa Szkoła Zarządzania
i Bankowości w Krakowie

Wyższe Szkoły Bankowe

Media Partners



Made in Wro
Do IT here!





Introducing SQLProxy

Maciej Pilecki

Microsoft Certified Master



DATA MASTER
LET IT SCALE




About me...

- Microsoft Certified Master in SQL Server
- I've been a SQL Server consultant for quite a while...
- Working on large-scale SQL Server deployments worldwide
 - Large-scale is not the same as big-data
 - It's the workload that matters most
- I've seen my share of scalability problems


THE PROBLEM WITH DATABASE SCALABILITY

Everything changes at large scale

- As the workload increases and becomes more complex, you face more and more contention on various resources:
 - Locks
 - Latches
 - Spinlocks
 - Memory
 - Tempdb
 - Even the Query Optimizer itself




What it MIGHT look like




Two approaches to scalability

- Scale-up – „let’s just buy a bigger server”
 - Expensive (Big Iron, Enterprise Edition licensing)
 - Does not always work (esp. for contention related problems)
- Scale-out – „let’s spread it across many smaller servers”
 - Usually, some form of data distribution and workload offloading
 - Replication, Federated Databases, Log Shipping, AlwaysOn Readable Secondaries
 - Complicated to the point of impossibility
 - Requires code changes and smart application design
 - AlwaysOn Readable Secondaries are cool, but...




The problem with AlwaysOn

- Readable Secondaries allow R/O queries
- Sounds like a perfect solution for scale-out and load-balancing, but...
- There's nothing in the SQL Native Client to enable that
- In order to do that in practice you need:
 - Two database connections in the application
 - Some logic in the application to distinguish R/W from R/O
 - All that means serious code and architecture changes
 - Most likely not possible for an existing, deployed application




Interesting thing about workloads...

- In most workloads, vast majority of batches are SELECTs
- These are often small, R/O queries, looking quite innocent...
- But mixed with even a small number of writes, they become a WMD
- Reads blocking writes
- Writes blocking reads
- Lots and lots of contention...



Another interesting thing...

- Many workloads include a high number of queries that return the same data each time they are executed
- The application asks the same question, over and over and over...
- And gets the same answer...
- Even if the data changes, it changes infrequently
- But the work has to be done each time:
 - Resource usage (CPU, memory, tempdb etc.)
 - Possible contention in many places
- That means lots of „unnecessary” queries



I STARTED LOOKING FOR A SIMPLE SOLUTION...

I asked myself some questions...

- What if we could cache queries?
- What if we could scale-out without the need for code changes?
- What if we could fix some of those „monster queries“?
- Wouldn't that be a dream-tool of every SQL Server consultant?


Introducing SQLProxy




- Fully transparent TDS proxy
- SQL Server and T-SQL aware
- Functionalities include:
 - Routing of R/O queries to the Readable Secondary
 - Caching of the query results
 - Dynamic query rewrites

Routing of read-only queries

- Allows scale-out and load-balancing using AlwaysOn Readable Secondaries
- Currently, SNAC does not support this natively and major application changes are needed to achieve that
- SQLProxy provides query routing that is:
 - fully transparent to the application
 - fully configurable
 - customizable




Routing of read-only queries




How R/O routing works

- Configure SQLProxy to connect to the AW Listener
- SQLProxy uses AW routing information to open a R/O connection to the Readable Secondary Replica
- Application opens a single connection to the SQLProxy
- Every batch is classified and executed on an appropriate server
- Classification can be:
 - Declarative (you specify which queries should be routed where)
 - Dynamic (by parsing the batch text and examining it)




Query result caching

- SQL Server caches a lot of things, including:
 - Data pages
 - Execution plans
- But your query is executed EVERY TIME you send it, taking server resources
- Solution:
 - Caching the query results




Data caching

First Execution




Subsequent Executions




How query caching works

- Fully transparent to the application
- Fully configurable, with cache expiration
- Works for R/O queries only (obviously)
- Works with parameterized queries and SP calls
- Time-bases expiration
- Configurable „on the fly”, with management interface




What about cache invalidation?

- Forget about it! 😊
- Too complex and resource intensive to implement in practice
- We use pure time-based expiration
- TTL of the result in cache is your decision:
 - Some applications will be just fine with the „stale” data
 - Some other may require more up-to-date information




What about scaling the SQLProxy itself?

- Will it become a bottleneck or single point of failure?
- NO!
- It's just a gateway and it's more-or-less stateless, so no problem:
- Just have more instances of SQLProxy and spread the load
- SQLProxy supports Windows Failover Clustering
- And it works with the –FailoverPartner switch in SNAC



Scaling the SQLProxy




What's in version 1.0?

- Query result caching
- Query routing between AG replicas
- Early implementation of on-the-fly query transformations
- Server connection pooling
- With some limitations:
 - SQL Authentication only in v. 1.0

What more can we do with it...


- Some ideas for version 2.0:
 - Auditing, logging etc.
 - Security (enforcing access rules, preventing SQL injection)
 - Error detection/correction (i.e. 1205 – Deadlock victim)
 - Translation between SQL dialects
 - Application compatibility with new versions of SQL Server
 - Data obfuscation
 - T-SQL language extensions



What's next?



- We are in a private Beta stage
- If you are interested, we want to talk to you!
- Wanna try it with your application? We want to talk to you!!!
- Come and grab me during the Party, I will answer all your questions
- AND we can have a drink together! 😊
- Visit us at www.sqlproxy.net for updates and more info



www.sqlproxy.net

THANK YOU!



DATA MASTER
LET IT SCALE



Strategic Sponsor



Gold Sponsors



Silver Sponsors



Technical Partners



Academic Partners



Wyższa Szkoła Zarządzania
i Bankowości w Krakowie

Wyższe Szkoły Bankowe

Media Partners



Made in Wro
Do IT here!

